

# califoenia

October 2, 2024

```
[20]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[21]: df = pd.read_csv('housing.csv')
df
```

```
[21]: longitude latitude housing_median_age total_rooms total_bedrooms \
0      -122.23    37.88           41.0         880.0         129.0
1      -122.22    37.86           21.0        7099.0        1106.0
2      -122.24    37.85           52.0        1467.0         190.0
3      -122.25    37.85           52.0        1274.0         235.0
4      -122.25    37.85           52.0        1627.0         280.0
...      ...      ...           ...           ...           ...
20635   -121.09    39.48           25.0        1665.0         374.0
20636   -121.21    39.49           18.0         697.0         150.0
20637   -121.22    39.43           17.0        2254.0         485.0
20638   -121.32    39.43           18.0        1860.0         409.0
20639   -121.24    39.37           16.0        2785.0         616.0

population households median_income
median_house_value \
0          322.0       126.0        8.3252        452600.0
1        2401.0       1138.0        8.3014        358500.0
2          496.0       177.0        7.2574        352100.0
3          558.0       219.0        5.6431        341300.0
4          565.0       259.0        3.8462        342200.0
...      ...      ...           ...           ...
20635        845.0       330.0        1.5603         78100.0
20636        356.0       114.0        2.5568         77100.0
20637       1007.0       433.0        1.7000         92300.0
20638        741.0       349.0        1.8672         84700.0
```

```

20639      1387.0      530.0      2.3886      89400.0

      ocean_proximity
0          NEAR BAY
1          NEAR BAY
2          NEAR BAY
3          NEAR BAY
4          NEAR BAY
...
20635          INLAND
20636          INLAND
20637          INLAND
20638          INLAND
20639          INLAND

[20640 rows x 10 columns]

```

## 0.1 Task 1: Identifying Variable Types

```
[22]: df.dtypes
```

```

[22]: longitude      float64
latitude          float64
housing_median_age float64
total_rooms       float64
total_bedrooms    float64
population        float64
households        float64
median_income     float64
median_house_value float64
ocean_proximity   object
dtype: object

```

```
[23]: df.info()
```

```

<class
'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to
20639 Data columns (total 10
columns):
#   Column                Non-Null Count  Dtype
---  -
0 longitude             20640 non-null float64
1 latitude              20640 non-null float64
2 housing_median_age    20640 non-null float64
3 total_rooms           20640 non-null float64

```

```

4  total_bedrooms      20433      non-null
                                float64
5  population          20640      non-null
                                float64
6  households          20640      non-null
                                float64
7  median_income       20640      non-null
                                float64
8    median_house_value 20640 non-null float64
9    ocean_proximity    20640 non-null object dtypes:
float64(9), object(1)
memory usage: 1.6+ MB

```

```
[24]: df.isnull().sum().sort_values(ascending=False)
```

```

[24]: total_bedrooms      207
longitude                0
latitude                 0
housing_median_age       0
total_rooms              0
population               0
households               0
median_income            0
median_house_value       0
ocean_proximity          0
dtype: int64

```

```
[25]: numerical_columns = df.select_dtypes(np.number)
numerical_columns.sample(6)
```

```

[25]:   longitude  latitude  housing_median_age  total_rooms  total_bedrooms \
4535   -118.21    34.03             44.0      1550.0         407.0
12750  -121.38    38.62             34.0      2352.0         610.0
8327   -118.30    33.94             36.0      2041.0         531.0
2709   -115.69    32.79             18.0      1564.0         340.0
15813  -122.42    37.76             52.0      4407.0        1192.0
7720   -118.11    33.94             37.0      1434.0         262.0

      population  households  median_income
median_house_value
4535      1718.0      403.0      2.5268      141100.0
12750      1127.0      592.0      2.2000      116500.0
8327      1390.0      464.0      2.0114      99300.0
2709      1161.0      343.0      2.1792      55200.0
15813      2280.0     1076.0      3.3937     270000.0

```

7720          786.0          256.0          4.4375          244900.0

```
[26]: categoric = df.select_dtypes(object)
      categoric.sample(5)
```

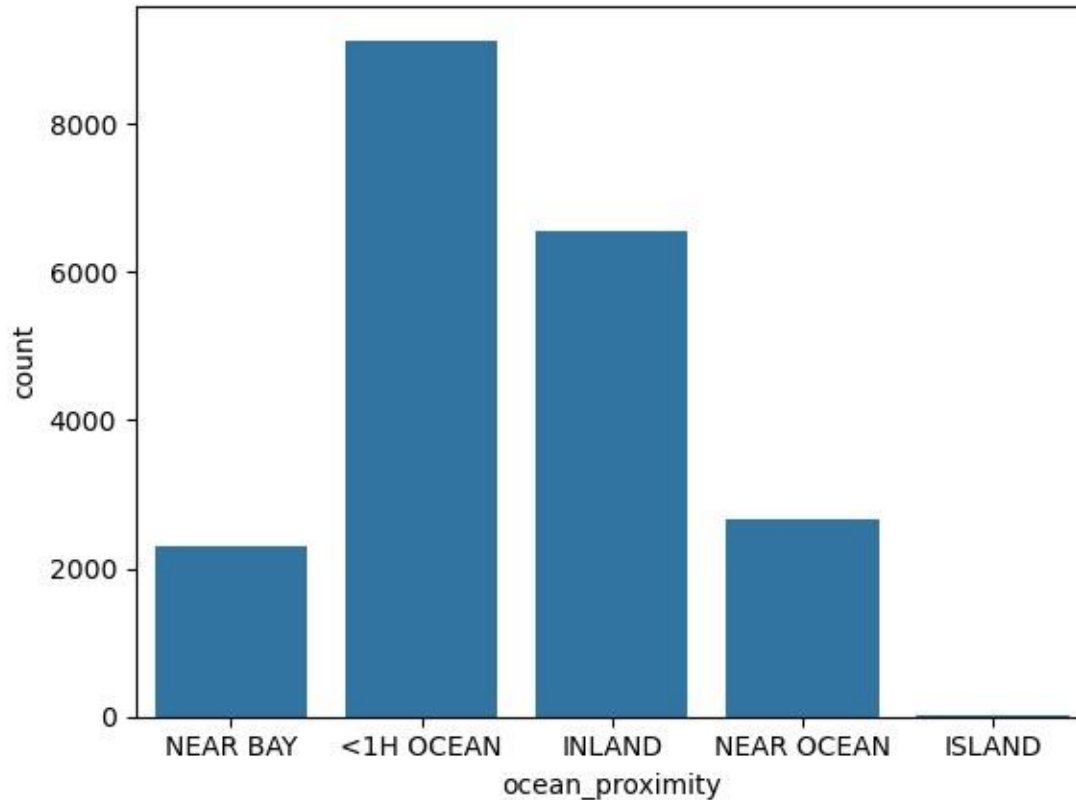
```
[26]: ocean_proximity
      2416          INLAND
      4930 <1H OCEAN 20229
      NEAR OCEAN
      5250 <1H OCEAN 8114
      NEAR OCEAN
```

```
[27]: categoric.value_counts()
```

```
[27]: ocean_proximity
      <1H OCEAN          9136
      INLAND            6551
      NEAR OCEAN        2658
      NEAR BAY          2290
      ISLAND            5
      Name: count, dtype: int64
```

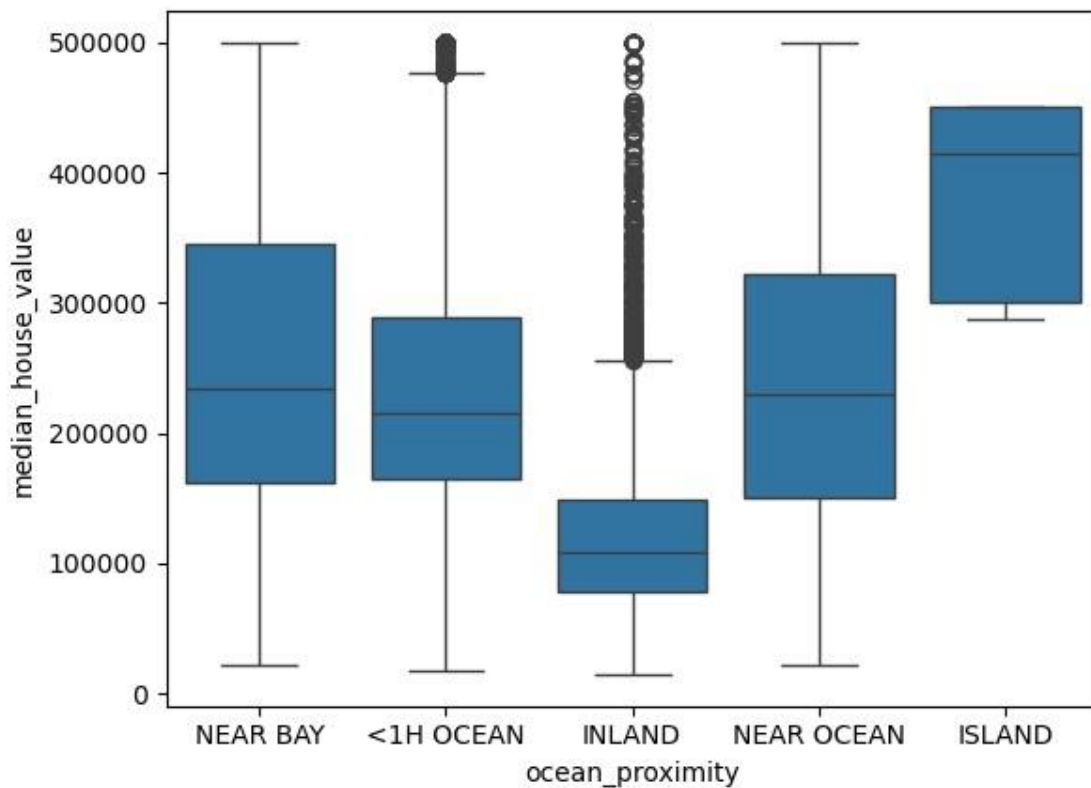
```
[28]: sns.countplot(x='ocean_proximity', data=df)
```

```
[28]: <Axes: xlabel='ocean_proximity', ylabel='count'>
```



```
[29]: sns.boxplot(x='ocean_proximity', y='median_house_value', data=df)
```

```
[29]: <Axes: xlabel='ocean_proximity', ylabel='median_house_value'>
```



## 0.2 Feature Scaling

### 0.2.1 Standardization

Standardization transforms the features to have a mean of 0 and a standard deviation of 1.

```
[30]:
```

```
scaler = StandardScaler()
scaled_data = scaler.fit_transform(numerical_columns)

standardized_df = pd.DataFrame(scaled_data, columns=numerical_columns.columns)
```

```
[31]: standardized_df
```

```
from sklearn.preprocessing import StandardScaler
```

```
[31]: longitude latitude housing_median_age total_rooms total_bedrooms \
```

```

0      -1.327835  1.052548          0.982143  -0.804819      -
                                           0.970325
1      -1.322844  1.043185          -0.607019  2.045890        1.348276
2      -1.332827  1.038503          1.856182  -0.535746      -
                                           0.825561
3      -1.337818  1.038503          1.856182  -0.624215      -
                                           0.718768
4      -1.337818  1.038503          1.856182  -0.462404      -
                                           0.611974
...      ...      ...      ...      ...      ...
20635  -0.758826  1.801647          -0.289187  -0.444985      -
                                           0.388895
20636  -0.818722  1.806329          -0.845393  -0.888704      -
                                           0.920488
20637  -0.823713  1.778237          -0.924851  -0.174995      -
                                           0.125472
20638  -0.873626  1.778237          -0.845393  -0.355600      -
                                           0.305834
20639  -0.833696  1.750146          -1.004309   0.068408        0.185416

      population households median_income
      median_house_value
0      -0.974429      -0.977033  2.344766  2.129631
1      0.861439  1.669961  2.332238  1.314156
2      -0.820777      -0.843637  1.782699  1.258693
3      -0.766028      -0.733781  0.932968  1.165100
4      -0.759847      -0.629157  -0.012881  1.172900
...      ...      ...      ...      ...
20635  -0.512592 -0.443449  -1.216128  -1.115804
20636  -0.944405 -1.008420  -0.691593  -1.124470
20637  -0.369537 -0.174042  -1.142593  -0.992746
20638  -0.604429 -0.393753  -1.054583  -1.058608
20639  -0.033977 0.079672  -0.780129  -1.017878
[20640 rows x 9 columns]

```

## 0.2.2 Normalization

Normalization scales the data between 0 and 1. This is useful when the scale of features differs, and you want to maintain relative differences.

```
[32]: from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()

scaled_data = scaler.fit_transform(numerical_columns)

normalized_df = pd.DataFrame(scaled_data, columns=numerical_columns.columns)

normalized_df
```

```
[32]: longitude latitude housing_median_age total_rooms total_bedrooms \
0      0.211155 0.567481      0.784314      0.022331      0.019863
1      0.212151 0.565356      0.392157      0.180503      0.171477
2      0.210159 0.564293      1.000000      0.037260      0.029330
3      0.209163 0.564293      1.000000      0.032352      0.036313
4      0.209163 0.564293      1.000000      0.041330      0.043296
...      ...      ...      ...      ...      ...
20635  0.324701 0.737513      0.470588      0.042296      0.057883
20636  0.312749 0.738576      0.333333      0.017676      0.023122
20637  0.311753 0.732200      0.313725      0.057277      0.075109
20638  0.301793 0.732200      0.333333      0.047256      0.063315
20639  0.309761 0.725824      0.294118      0.070782      0.095438


      population households median_income median_house_value
0      0.008941  0.020556  0.539668      0.902266
1      0.067210  0.186976  0.538027      0.708247
2      0.013818  0.028943  0.466028      0.695051
3      0.015555  0.035849  0.354699      0.672783
4      0.015752  0.042427  0.230776      0.674638
...      ...      ...      ...      ...
20635  0.023599  0.054103  0.073130      0.130105
20636  0.009894  0.018582  0.141853      0.128043
20637  0.028140  0.071041  0.082764      0.159383
20638  0.020684  0.057227  0.094295      0.143713
```

```
20639    0.038790    0.086992    0.130253    0.153403
[20640 rows x 9 columns]
```

### 0.2.3 Impact of Scaling

```
[33]: # Before Scaling
sns.scatterplot(data=numerical_columns)
```

[33]: <Axes: >

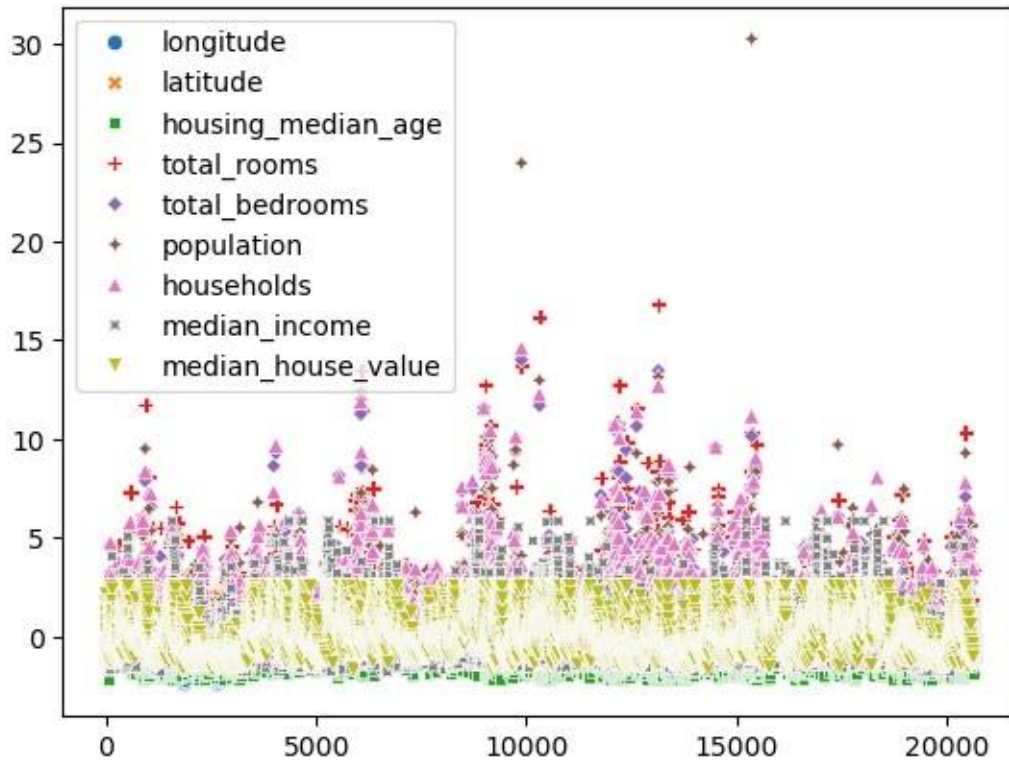


```
[34]: # After Scaling (Standardization)
sns.scatterplot(data=standardized_df)
```

[34]: <Axes: >


```
c:\Users\Abuzar\miniconda3\envs\ML\lib\sitepackages\IPython\core\pylab
tools.py:152: UserWarning: Creating legend with loc="best" can be slow
with large amounts of data. fig.canvas.print_figure(bytes_io, **kw)
```





```
[35]: # After Normalization
sns.scatterplot(data=normalized_df)
```

```
[35]: <Axes: >
```



```
[ ]:
```