

```

import numpy as np
import matplotlib.pyplot as plt

# Set random seed for reproducibility
np.random.seed(0)

# Generate synthetic dataset with features on different scales
X_small = np.random.rand(100, 1)           # Small scale
feature (0-1)
X_large = np.random.rand(100, 1) * 1000    # Large scale
feature (0-1000)
X = np.hstack((X_small, X_large))          # Combine features

# True coefficients for linear regression (for generating target values)
true_coef = np.array([5, 0.01])
y = X @ true_coef + np.random.randn(100) * 5 # Generate target variable with noise

# Gradient Descent function
def gradient_descent(X, y, learning_rate, iterations):
    n_samples = X.shape[0]
    theta = np.zeros(X.shape[1])
    cost_history = []
    for i in range(iterations):
        y_pred = X @ theta
        error = y_pred - y
        theta -= (learning_rate / n_samples) * (X.T @ error)
    cost = (1 / (2 * n_samples)) * np.sum(error ** 2)
    cost_history.append(cost)
    return theta, cost_history

# Run gradient descent on unscaled data with a small learning rate
theta_unscaled, cost_history_unscaled = gradient_descent(X, y,
learning_rate=0.0000001, iterations=100)

# Run gradient descent on the small dataset
theta_small, cost_history_small = gradient_descent(X_small, y,
learning_rate=0.0000001, iterations=100)

# Run gradient descent on the large dataset
theta_large, cost_history_large = gradient_descent(X_large, y,
learning_rate=0.0000001, iterations=100)

# Plot cost history for both datasets
plt.figure(figsize=(12, 6)) # Plot
for small dataset

```

```

plt.subplot(1, 2, 1)
plt.plot(cost_history_small, color="blue")
plt.title("Gradient Descent on Small Scale Feature")
plt.xlabel("Iterations") plt.ylabel("Cost (MSE)")
plt.grid(True)

# Plot for large dataset
plt.subplot(1, 2, 2)
plt.plot(cost_history_large, color="orange")
plt.title("Gradient Descent on Large Scale Feature")
plt.xlabel("Iterations") plt.ylabel("Cost (MSE)")
plt.grid(True)

plt.tight_layout()
plt.show()

```

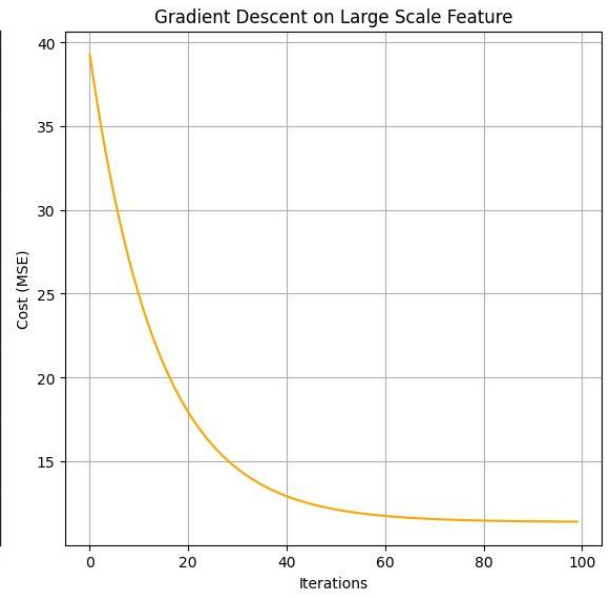
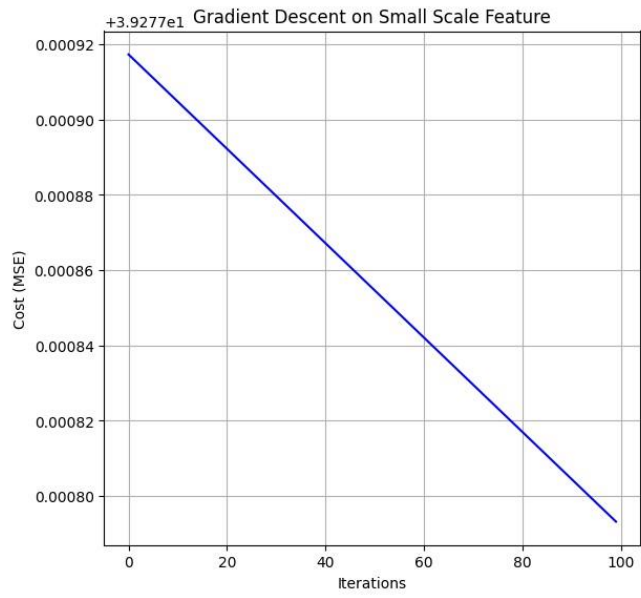
```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler

# Set random seed for reproducibility
np.random.seed(0)

# Generate synthetic dataset with features on different scales
X_small = np.random.rand(100, 1) # Small scale
feature (0-1)
X_large = np.random.rand(100, 1) * 1000 # Large scale
feature (0-1000)
X = np.hstack((X_small, X_large)) # Combine features

```



```

# True coefficients for linear regression (for generating target
values)
true_coef = np.array([5, 0.01])
y = X @ true_coef + np.random.randn(100) * 5          # Generate target
variable with noise

# Apply Min-Max scaling to the entire dataset
scaler_minmax = MinMaxScaler()
X_minmax = scaler_minmax.fit_transform(X)

# Apply Min-Max scaling separately to small and large features
scaler_small = MinMaxScaler()
X_small_minmax = scaler_small.fit_transform(X_small)

scaler_large = MinMaxScaler()
X_large_minmax = scaler_large.fit_transform(X_large)

# Gradient Descent function
def gradient_descent(X, y,
learning_rate, iterations):    n_samples = X.shape[0]
theta = np.zeros(X.shape[1])    cost_history = []
for i in range(iterations):
    y_pred = X @ theta
    error = y_pred - y
    theta -= (learning_rate / n_samples) * (X.T @ error)
    cost = (1 / (2 * n_samples)) * np.sum(error ** 2)
    cost_history.append(cost)    return theta, cost_history

# Run gradient descent on Min-Max scaled data
theta_minmax, cost_history_minmax = gradient_descent(X_minmax, y,
learning_rate=0.01, iterations=100)
theta_small_minmax, cost_history_small_minmax =
gradient_descent(X_small_minmax, y, learning_rate=0.01,
iterations=100)
theta_large_minmax, cost_history_large_minmax =
gradient_descent(X_large_minmax, y, learning_rate=0.01,
iterations=100)

# Plot cost history for all datasets
plt.figure(figsize=(15, 6))

# Plot for all Min-Max scaled data
plt.subplot(1, 3, 1)
plt.plot(cost_history_minmax, color="green")
plt.title("Gradient Descent on Min-Max Scaled Data")
plt.xlabel("Iterations") plt.ylabel("Cost (MSE)")

```

```

plt.grid(True)

# Plot for Min-Max scaled small dataset
plt.subplot(1, 3, 2)
plt.plot(cost_history_small_minmax, color="blue")
plt.title("Gradient Descent on Small Scale Min-Max Data")
plt.xlabel("Iterations") plt.ylabel("Cost (MSE)")
plt.grid(True)

# Plot for Min-Max scaled large dataset
plt.subplot(1, 3, 3)
plt.plot(cost_history_large_minmax, color="orange")
plt.title("Gradient Descent on Large Scale Min-Max Data")
plt.xlabel("Iterations") plt.ylabel("Cost (MSE)")
plt.grid(True)

plt.tight_layout()
plt.show()

```




```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler

# Set random seed for reproducibility
np.random.seed(0)

# Generate synthetic dataset with features on different scales
X_small = np.random.rand(100, 1) # Small scale
feature (0-1)
X_large = np.random.rand(100, 1) * 1000 # Large scale
feature (0-1000)
X = np.hstack((X_small, X_large)) # Combine features

```



```

# True coefficients for linear regression (for generating target
values)
true_coef = np.array([5, 0.01])
y = X @ true_coef + np.random.randn(100) * 5          # Generate target
variable with noise

# Apply Standardization to the entire dataset
scaler_standard = StandardScaler()
X_standard = scaler_standard.fit_transform(X)

# Apply Standardization separately to small and large features
scaler_small = StandardScaler()
X_small_standard = scaler_small.fit_transform(X_small)

scaler_large = StandardScaler()
X_large_standard = scaler_large.fit_transform(X_large)

# Gradient Descent function
def gradient_descent(X, y,
learning_rate, iterations):    n_samples = X.shape[0]
theta = np.zeros(X.shape[1])    cost_history = []
for i in range(iterations):
    y_pred = X @ theta
    error = y_pred - y
    theta -= (learning_rate / n_samples) * (X.T @ error)
cost = (1 / (2 * n_samples)) * np.sum(error ** 2)
cost_history.append(cost)    return theta, cost_history

# Run gradient descent on standardized data
theta_standard, cost_history_standard = gradient_descent(X_standard,
y, learning_rate=0.01, iterations=100)
theta_small_standard, cost_history_small_standard =
gradient_descent(X_small_standard, y, learning_rate=0.01,
iterations=100)
theta_large_standard, cost_history_large_standard =
gradient_descent(X_large_standard, y, learning_rate=0.01,
iterations=100)

# Plot cost history for all datasets
plt.figure(figsize=(15, 6))

# Plot for standardized entire dataset
plt.subplot(1, 3, 1)
plt.plot(cost_history_standard, color="red")
plt.title("Gradient Descent on Standardized Data")
plt.xlabel("Iterations") plt.ylabel("Cost (MSE)")

```

```

plt.grid(True)

# Plot for standardized small dataset
plt.subplot(1, 3, 2)
plt.plot(cost_history_small_standard, color="blue")
plt.title("Gradient Descent on Small Scale Standardized Data")
plt.xlabel("Iterations") plt.ylabel("Cost (MSE)")
plt.grid(True)

# Plot for standardized large dataset
plt.subplot(1, 3, 3)
plt.plot(cost_history_large_standard, color="orange")
plt.title("Gradient Descent on Large Scale Standardized Data")
plt.xlabel("Iterations") plt.ylabel("Cost (MSE)")
plt.grid(True)

plt.tight_layout()
plt.show()

```




```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler

# Set random seed for reproducibility
np.random.seed(0)

# Generate synthetic dataset with features on different scales
X_small = np.random.rand(100, 1) # Small scale
feature (0-1)
X_large = np.random.rand(100, 1) * 1000 # Large scale
feature (0-1000)
X = np.hstack((X_small, X_large)) # Combine features

```

```

# True coefficients for linear regression (for generating target
values)
true_coef = np.array([5, 0.01])
y = X @ true_coef + np.random.randn(100) * 5      # Generate target
variable with noise

# Apply Standardization to the entire dataset
scaler_standard = StandardScaler()
X_standard = scaler_standard.fit_transform(X)

# Apply Standardization separately to small and large features
scaler_small = StandardScaler()
X_small_standard = scaler_small.fit_transform(X_small)

scaler_large = StandardScaler()
X_large_standard = scaler_large.fit_transform(X_large)

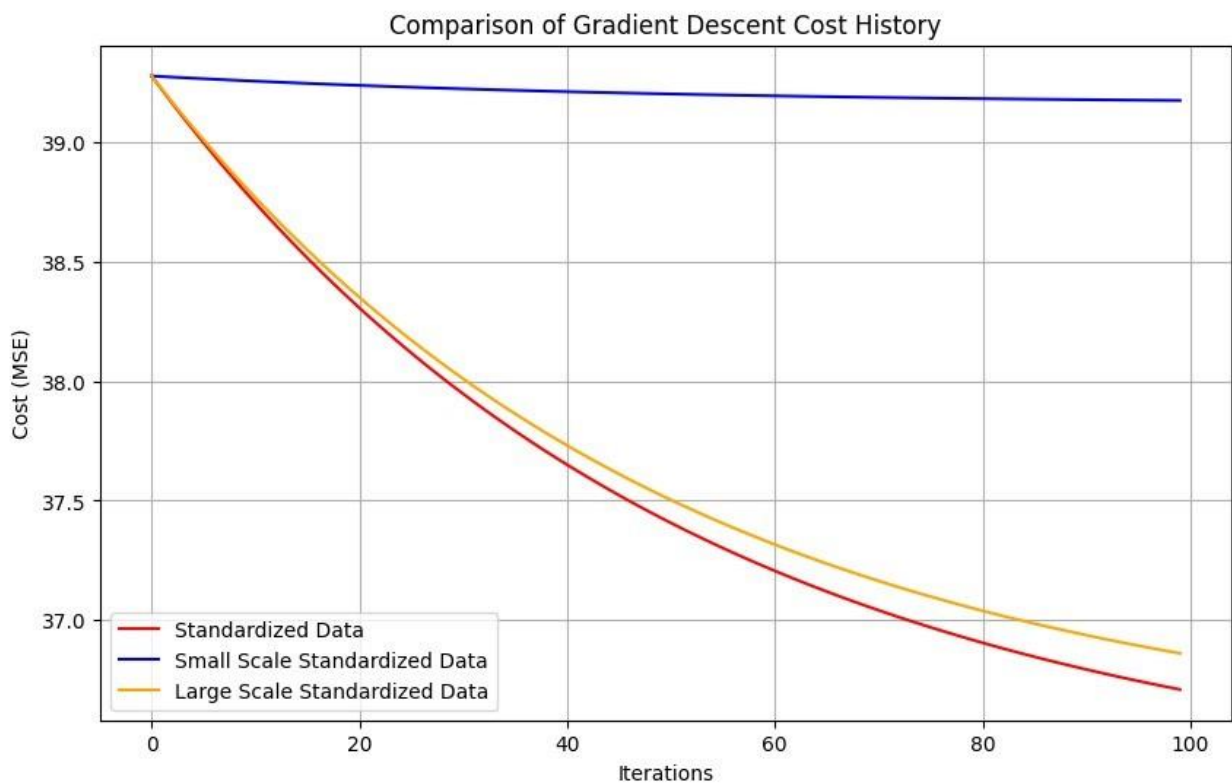
# Gradient Descent function
def gradient_descent(X, y,
learning_rate, iterations):    n_samples = X.shape[0]
theta = np.zeros(X.shape[1])    cost_history = []
for i in range(iterations):
    y_pred = X @ theta
error = y_pred - y
    theta -= (learning_rate / n_samples) * (X.T @ error)
cost = (1 / (2 * n_samples)) * np.sum(error ** 2)
cost_history.append(cost)    return theta, cost_history

# Run gradient descent on standardized data
theta_standard, cost_history_standard = gradient_descent(X_standard,
y, learning_rate=0.01, iterations=100)
theta_small_standard, cost_history_small_standard =
gradient_descent(X_small_standard, y, learning_rate=0.01,
iterations=100)
theta_large_standard, cost_history_large_standard =
gradient_descent(X_large_standard, y, learning_rate=0.01,
iterations=100)

# Plot cost history for comparison
plt.figure(figsize=(10, 6))
plt.plot(cost_history_standard, label="Standardized Data",
color="red")
plt.plot(cost_history_small_standard, label="Small Scale Standardized
Data", color="blue")
plt.plot(cost_history_large_standard, label="Large Scale Standardized
Data", color="orange")

```

```
plt.title("Comparison of Gradient Descent Cost History")
plt.xlabel("Iterations") plt.ylabel("Cost (MSE)")
plt.legend() plt.grid(True) plt.show()
```



```
print("""
Explanation of Results:
Unscaled Data:

Gradient descent struggles to converge due to the large differences in
feature ranges (0-1 vs. 0-1000), and thus requires a tiny learning
rate. Convergence is very slow or may not happen in a reasonable
timeframe.
Min-Max Scaling:

Scales all features to the
[
0
,
1
]
[0,1] range, allowing gradient descent to converge steadily at a
reasonable learning rate.

This scaling method works well but can be sensitive to extreme values.
Standardization:
```

Standardization (mean=0, std=1) scales features more effectively when large variances are present, stabilizing updates even with outliers. It tends to be more robust than Min-Max scaling and is generally the preferred choice for gradient descent.""")

Explanation of Results:

Unscaled Data:

Gradient descent struggles to converge due to the large differences in feature ranges (0-1 vs. 0-1000), and thus requires a tiny learning rate. Convergence is very slow or may not happen in a reasonable timeframe.

Min-Max Scaling:

Scales all features to the

[
0
,
1
]

[0,1] range, allowing gradient descent to converge steadily at a reasonable learning rate. This scaling method works well but can be sensitive to extreme values.

Standardization:

Standardization (mean=0, std=1) scales features more effectively when large variances are present, stabilizing updates even with outliers. It tends to be more robust than Min-Max scaling and is generally the preferred choice for gradient descent.