

Lab_Task06

October 16, 2024

```
[9]: import pandas as pd
df=pd.read_csv( C:/5th_semester/machine_learning/ml/data.csv )
```

```
[3]: df
```

```
[3]:
```

	TV	Radio	Newspaper	Sales
0	143.635030	35.681266	7.578097	14.015299
1	287.678577	13.365599	27.563823	18.227685
2	232.998485	16.465149	17.631309	16.492024
3	199.664621	45.942168	25.661437	19.667324
4	89.004660	34.257162	13.001240	11.494491
..
195	137.302394	47.230293	16.849041	19.388495

[200 rows x 4 columns]

```
[4]: import seaborn as sns
import matplotlib.pyplot as plt

# Set the aesthetics for the plots
sns.set(style="whitegrid")

# 1. Pairplot - To visualize pairwise relationships between features and target □
□variable
plt.figure(figsize=(10 6))
sns.pairplot(data)
plt.suptitle("Pairplot of Advertising Data    =1.02)
plt.show()

# 2. Correlation Heatmap - To visualize the correlations between features and □
□target
plt.figure(figsize=(8 6))
corr_matrix = data.corr()

196 231.488920 44.336510 21.688943 26.308555
197 274.277565 27.159761 9.307997 22.329760
198 271.771606 40.034843 9.807225 22.038494
199 244.968886 40.181715 6.021715 18.673907
```

```

sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', linewidths=0.5,
            square=True)
plt.title("Correlation Heatmap")
plt.show()

# 3. Histograms - To visualize the distribution of individual features
data.hist(bins=15, figsize=(10, 8), edgecolor='black')
plt.suptitle("Histograms of Features and Target")
plt.show()

# 4. Scatter Plots - To visualize relationships between features and the target
            (Sales)
plt.figure(figsize=(15, 5))

# TV vs Sales
plt.subplot(1, 3, 1)
plt.scatter(data['TV'], data['Sales'], color='blue', alpha=0.5)
plt.title("TV vs Sales")
plt.xlabel("TV Advertising Spend")
plt.ylabel("Sales")

# Radio vs Sales
plt.subplot(1, 3, 2)
plt.scatter(data['Radio'], data['Sales'], color='green', alpha=0.5)
plt.title("Radio vs Sales")
plt.xlabel("Radio Advertising Spend")
plt.ylabel("Sales")

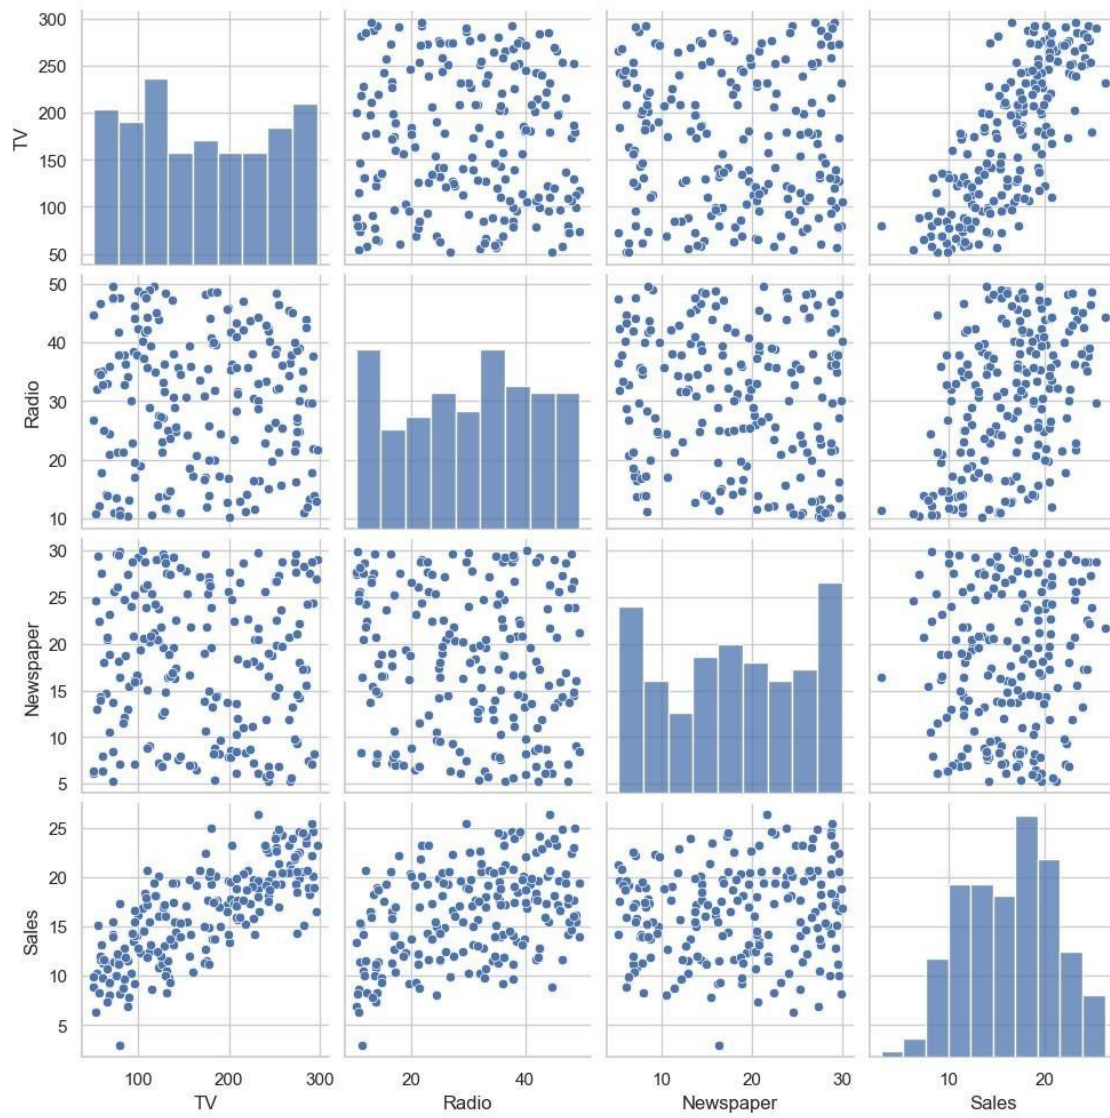
# Newspaper vs Sales
plt.subplot(1, 3, 3)
plt.scatter(data['Newspaper'], data['Sales'], color='red', alpha=0.5)
plt.title("Newspaper vs Sales")
plt.xlabel("Newspaper Advertising Spend")
plt.ylabel("Sales")

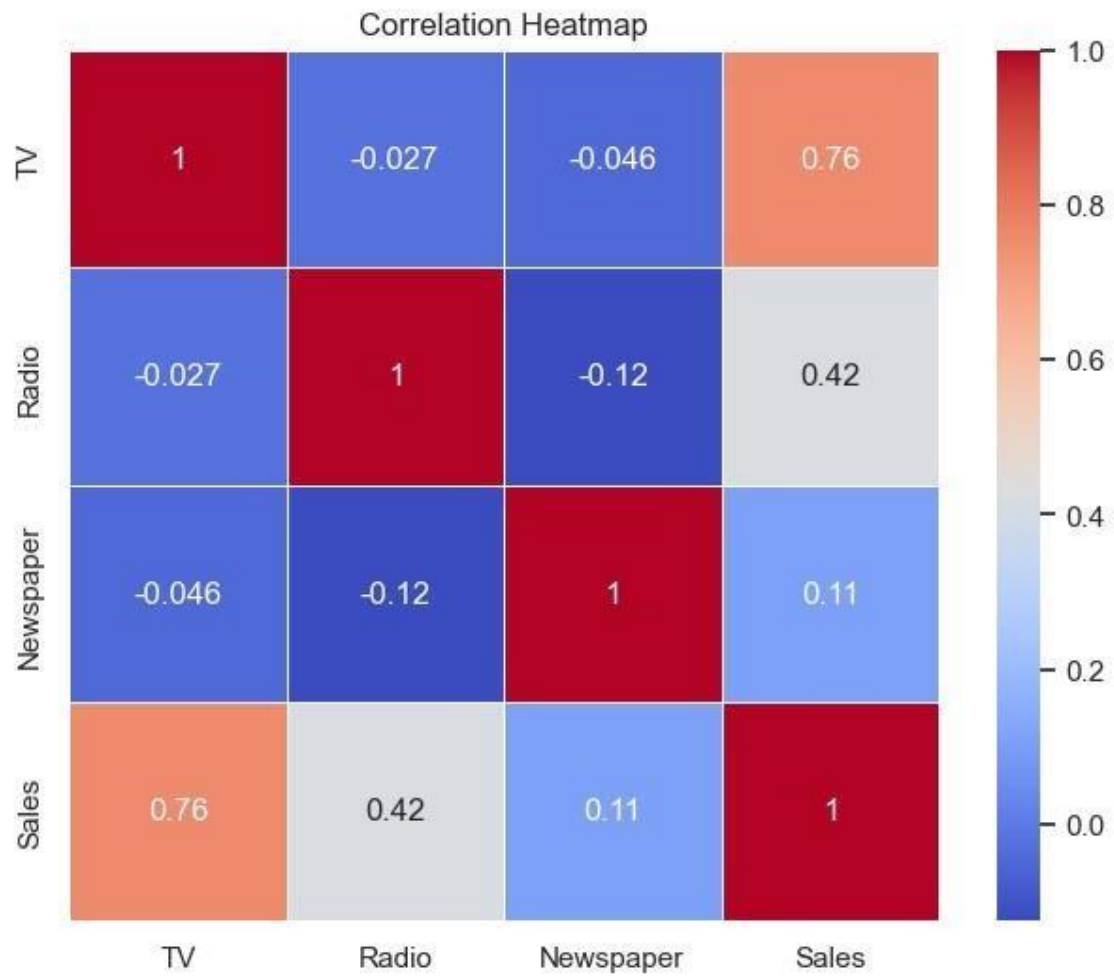
plt.tight_layout()
plt.show()

```

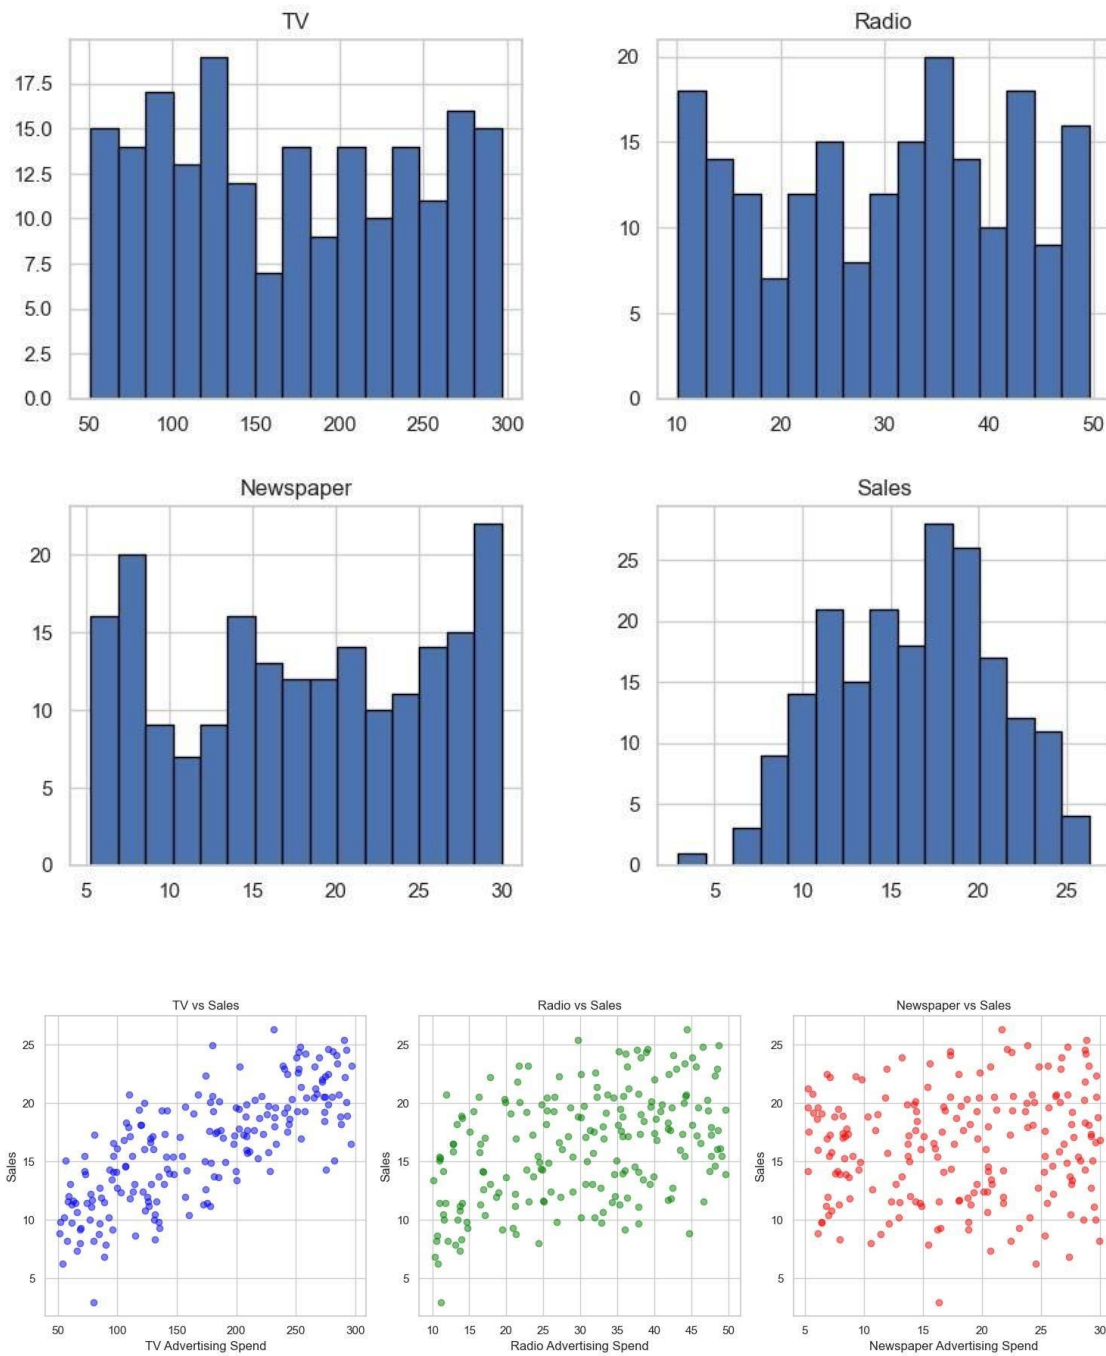
<Figure size 1000x600 with 0 Axes>

Pairplot of Advertising Data





Histograms of Features and Target



```
[5]: # Step 2: Implement Linear Regression from Scratch
```

```
class LinearRegressionScratch
```

```

def __init__(self):
    self.beta = None

def fit(self, X, y):
    # Add intercept (bias term) to X
    X_b = np.c_[np.ones((X.shape[0], 1)) , X] # Add a column of ones to X
    # Compute the weights using the normal equation
    self.beta = np.linalg.inv(X_b.T.dot(X_b)).dot(X_b.T).dot(y)

def predict(self, X):
    # Add intercept (bias term) to X
    X_b = np.c_[np.ones((X.shape[0], 1)) , X]
    # Make predictions
    return X_b.dot(self.beta)

# Splitting the data into training and testing sets
X = data[['TV', 'Radio', 'Newspaper']].values
y = data['Sales'].values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)

# Initialize and fit the model
model_scratch = LinearRegressionScratch()
model_scratch.fit(X_train, y_train)

# Predict on test data
y_pred_scratch = model_scratch.predict(X_test)

# Evaluate the model using Mean Squared Error (MSE)
mse_scratch = mean_squared_error(y_test, y_pred_scratch)
print(f"MSE from Scratch: {mse_scratch}")

```

MSE from Scratch: 2.7616556127332315

```

[6]: from sklearn.linear_model import LinearRegression

# Initialize the LinearRegression model from scikit-learn
model_sklearn = LinearRegression()

# Fit the model on the training data
model_sklearn.fit(X_train, y_train)

# Predict on test data
y_pred_sklearn = model_sklearn.predict(X_test)

# Evaluate the model using Mean Squared Error (MSE)

```

```
mse_sklearn = mean_squared_error(y_test, y_pred_sklearn)
print(f"MSE with scikit-learn:
{mse_sklearn}")
```

MSE with scikit-learn: 2.761655612733244

[7]: # Step 2: Calculate R-squared (accuracy) using
scikit-learn

```
r2_sklearn = model_sklearn.score(X_test, y_test)
print(f"R-squared with scikit-learn: {r2_sklearn}")
```

R-squared with scikit-learn: 0.8716759282807236

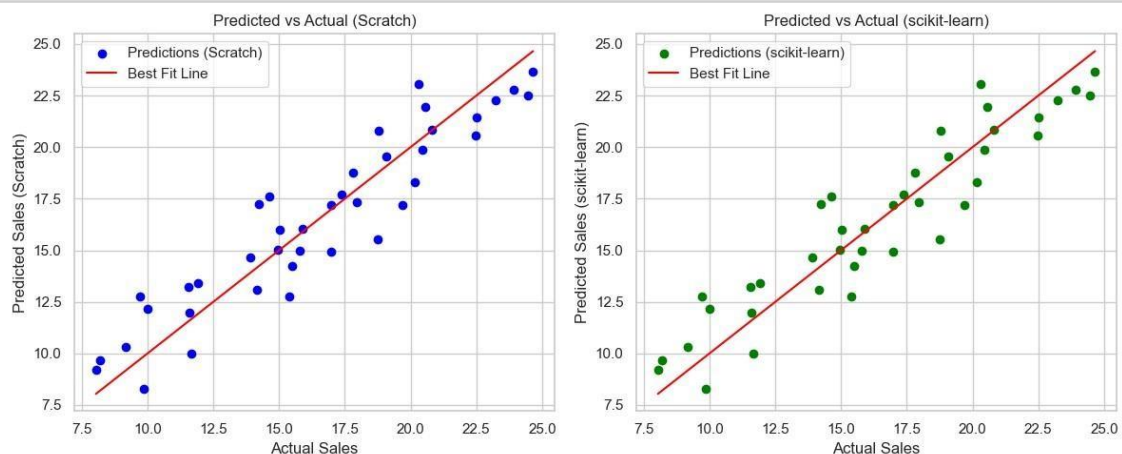

```
[8]: # Import required libraries
import matplotlib.pyplot as plt
import numpy as np

# Step 1: Plot predictions vs actual values
plt.figure(figsize=(12, 5))

# Predictions from scratch model
plt.subplot(1, 2, 1)
plt.scatter(y_test, y_pred_scratch, color='blue', label='Predictions (Scratch)')
# Plot best-fitted line
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         label='Best Fit Line')
plt.xlabel("Actual Sales")
plt.ylabel("Predicted Sales (Scratch)")
plt.title("Predicted vs Actual (Scratch)")
plt.legend()

# Predictions from scikit-learn model
plt.subplot(1, 2, 2)
plt.scatter(y_test, y_pred_sklern, color='green', label='Predictions (scikit-learn)')
# Plot best-fitted line
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         label='Best Fit Line')
plt.xlabel("Actual Sales")
plt.ylabel("Predicted Sales (scikit-learn)")
plt.title("Predicted vs Actual (scikit-learn)")
plt.legend()

plt.tight_layout()
plt.show()
```



[]: